

Foundations of Data Analysis & AI

Orientation, Data Taxonomy, Collection, and Quality

Course: Data Analysis & AI Fundamentals (Mediacrest Training College)

Instructor: Victor Wandera Lumumba, MSc.

Contents

1	Welcome to the Data & AI Ecosystem	3
1.1	Program Roadmap	3
1.2	The Capstone Expectation	3
2	The Core of Data Analysis	3
2.1	The DIKW Pyramid	3
2.2	The Data Analysis Lifecycle	4
2.3	Why Data Analysis Matters	4
3	Taxonomy of Data	4
3.1	Statistical Classification	4
3.2	Structural Classification	5
3.2.1	Structured Data	5
3.2.2	Unstructured Data	5
3.2.3	Semi-Structured Data	5
3.3	Common Data Formats	5
4	Modern Data Collection Methodologies	6
4.1	Primary vs. Secondary Data	6
4.2	Modern Acquisition Techniques	6
4.2.1	APIs (Application Programming Interfaces)	6
4.2.2	Web Scraping	6
4.2.3	IoT & Streaming Data	6
5	Data Quality & The Art of Preprocessing	6
5.1	The Golden Rule: GIGO	6
5.2	Dimensions of Data Quality	7
5.3	Handling Missing Data	7
5.4	Outlier Detection & Treatment	7
6	The Python Data Science Toolstack	8
6.1	The Core Libraries	8
7	Conclusion & Next Steps	8
7.1	Expectations for Success	8

1 Welcome to the Data & AI Ecosystem

Welcome to the **Data Analysis & AI Fundamentals Scholarship Program**, an intensive 6-week certification course designed in partnership with Mediacrest Training College. This program is not just about learning to code; it is about cultivating an **analytical mindset** to solve real-world problems in health, finance, agriculture, and beyond.

1.1 Program Roadmap

The curriculum is structured to take you from absolute fundamentals to deploying end-to-end machine learning solutions:

1. **Weeks 1-2: Foundations of Data Analysis.** Python programming, data structures, statistical thinking, and handling raw data.
2. **Weeks 3-4: Data Visualization & Essential Analytics.** Storytelling with data, advanced cleaning, feature engineering, and uncovering trends.
3. **Week 5: Machine Learning & AI Fundamentals.** Transitioning from statistical modeling to predictive algorithms (Scikit-Learn).
4. **Week 6: AI Applications, Ethics & Capstone.** Responsible AI, Generative AI (LLMs), and a comprehensive final project.

1.2 The Capstone Expectation

By Week 6, you will not just have theoretical knowledge. You will build an **end-to-end analytics project**. This means defining a problem, sourcing data, cleaning it, building a model, and presenting actionable insights—exactly what you will do in a professional data science role.

2 The Core of Data Analysis

Before we write a single line of Python, we must understand the terrain. Data is the new oil, but unrefined oil is useless. Data analysis is the refinery.

2.1 The DIKW Pyramid

To understand data analysis, one must understand the hierarchy of value:

- **Data:** Raw, unorganized facts, numbers, or symbols (e.g., 39.5).
- **Information:** Data with context and meaning (e.g., Patient body temperature is 39.5°C).
- **Knowledge:** Application of information; recognizing patterns (e.g., 39.5°C indicates a high fever, often linked to infection).
- **Wisdom:** Actionable insight based on knowledge (e.g., Administer antipyretics and run a malaria rapid diagnostic test).

Definition: Data Analysis

Data Analysis is the systematic process of inspecting, cleansing, transforming, and modeling data with the goal of discovering useful information, informing conclusions, and supporting decision-making. It is the bridge between raw *Data* and actionable *Wisdom*.

2.2 The Data Analysis Lifecycle

Professional data science follows a structured lifecycle. While various frameworks exist (like CRISP-DM or OSEMN), we will use a practical 6-step approach:

1. **Ask:** Define the business or research problem. What decision needs to be made?
2. **Collect:** Gather the necessary data from databases, APIs, or surveys.
3. **Clean:** Handle missing values, fix errors, and format data (often the most time-consuming step).
4. **Analyze:** Apply statistical tests or machine learning algorithms to find patterns.
5. **Visualize:** Create charts and dashboards to communicate findings to stakeholders.
6. **Act:** Implement the data-driven recommendations.

2.3 Why Data Analysis Matters

- **Evidence-Based Decision Making:** Moves organizations away from "gut feelings" to empirical strategies.
- **Operational Efficiency:** Identifies bottlenecks in supply chains or manufacturing.
- **Risk Management:** Predicts equipment failures (predictive maintenance) or financial defaults.
- **Personalization:** Powers the recommendation engines of Netflix, Spotify, and Amazon.

Real-World Case Study: Healthcare Analytics

Consider the **CardioPredict AI** and **Mental Distress Screening Apps**. By analyzing clinical and demographic data, machine learning models can identify patients at high risk for heart disease or psychological distress *before* severe symptoms manifest. This shifts healthcare from *reactive* (treating the sick) to *proactive* (maintaining health).

3 Taxonomy of Data

Not all data is created equal. The type of data you are working with dictates the statistical methods and machine learning algorithms you can use.

3.1 Statistical Classification

From a statistical perspective, data is divided into two main branches:

1. **Quantitative (Numerical):**
 - *Discrete:* Countable, usually integers (e.g., number of children, number of hospital visits).
 - *Continuous:* Measurable, can take any value within a range (e.g., height, weight, temperature, revenue).
2. **Qualitative (Categorical):**
 - *Nominal:* Categories without an inherent order (e.g., Blood Type, Country, Device Type).

- **Ordinal:** Categories with a meaningful order or ranking (e.g., Education Level, Customer Satisfaction ratings).

3.2 Structural Classification

From an IT and database perspective, data is classified by how it is organized:

3.2.1 Structured Data

Highly organized, typically in tabular format with rows (records) and columns (features).

- **Storage:** Relational Database Management Systems (RDBMS) like PostgreSQL, MySQL.
- **Analysis:** Easily queried using SQL; ideal for traditional statistical analysis.

3.2.2 Unstructured Data

Data without a predefined schema or organization. It is heavy and complex.

- **Examples:** Emails, social media posts, medical imaging (X-rays, MRIs), audio recordings, video.
- **Analysis:** Requires advanced techniques like Natural Language Processing (NLP) for text, or Computer Vision (CNNs) for images.

3.2.3 Semi-Structured Data

Does not conform to a strict tabular structure but contains tags or markers to separate elements.

- **Examples:** JSON (JavaScript Object Notation), XML, HTML, NoSQL databases (MongoDB).
- **Importance:** This is the native format of the web. When you pull data from an API, it almost always arrives as JSON.

3.3 Common Data Formats

- **CSV (Comma Separated Values):** The universal language of tabular data. Lightweight, but lacks support for complex data types or multiple sheets.
- **Parquet / Feather: Columnar storage formats.** Unlike CSV (which is row-based), Parquet stores data by column. This allows for massive compression and highly efficient querying when you only need to read a few columns from a massive dataset.
- **JSON:** Hierarchical, key-value pairs. Essential for web data and document stores.

Under the Hood: Why Columnar Storage (Parquet) Matters

If you have a dataset with 100 columns and 1 billion rows, and you only need the **age** column to calculate the average:

- **Row-based (CSV):** The computer must read every single row across all 100 columns into memory, then discard 99 columns. Extremely slow.
- **Columnar (Parquet):** The computer reads *only* the **age** column from the disk. It is exponentially faster and uses a fraction of the RAM.

4 Modern Data Collection Methodologies

Data collection is the foundation of your analytical pipeline. If your foundation is flawed, your models will fail.

4.1 Primary vs. Secondary Data

- **Primary Data:** Data you collect yourself specifically for your current research question. Examples: Clinical trials, custom surveys, IoT sensor deployments.
- **Secondary Data:** Data collected by someone else. Examples: Government census data, WHO health statistics, Kaggle datasets, historical financial records.

4.2 Modern Acquisition Techniques

4.2.1 APIs (Application Programming Interfaces)

APIs allow software systems to talk to each other. In data science, we use REST APIs to pull live, structured data from web services.

- **Example:** Pulling real-time weather data from OpenWeatherMap, or fetching tweets via the X (Twitter) API.
- **Format:** Data is typically returned in JSON format.

4.2.2 Web Scraping

Extracting data directly from the HTML structure of websites when an API is not available.

- **Tools:** Python's `BeautifulSoup` for static pages, `Selenium` or `Playwright` for dynamic, JavaScript-rendered pages.
- **Ethics:** Always check the `robots.txt` file of a website and respect rate limits to avoid overloading the host server.

4.2.3 IoT & Streaming Data

The Internet of Things (IoT) generates continuous, high-frequency time-series data.

- **Examples:** Smartwatches tracking heart rate, smart meters tracking electricity usage, telematics in fleet vehicles.
- **Challenge:** Requires specialized time-series analysis and streaming architectures (like Apache Kafka).

5 Data Quality & The Art of Preprocessing

In the real world, data is messy. It is generated by humans, faulty sensors, and legacy systems.

5.1 The Golden Rule: GIGO

Definition: Garbage In

A concept originating from computer science. If you feed inaccurate, biased, or nonsense data into a machine learning model, the model's predictions will be equally inaccurate, biased, or nonsensical. No algorithm can fix fundamentally bad data.

Fact: Data scientists spend **60% to 80%** of their time cleaning and preprocessing data.

5.2 Dimensions of Data Quality

To assess if data is "good", we evaluate it across several dimensions:

1. **Accuracy:** Does the data reflect reality? (e.g., A patient's age is recorded as 150).
2. **Completeness:** Are there missing values? (e.g., 30% of patients lack a blood pressure reading).
3. **Consistency:** Is data uniform across systems? (e.g., One database uses "Kenya", another uses "KE").
4. **Timeliness:** Is the data up to date? (e.g., Using 2010 census data for 2026 urban planning).
5. **Validity:** Does it conform to syntax rules? (e.g., Dates formatted as DD/MM/YYYY vs MM/DD/YYYY).
6. **Uniqueness:** Are there duplicate records?

5.3 Handling Missing Data

Missing data is not just a nuisance; the *reason* it is missing can contain valuable information. Statisticians classify missingness into three mechanisms (Rubin's classifications):

- **MCAR (Missing Completely at Random):** The missingness has no relationship with any data. (e.g., A lab sample was dropped by accident). *Solution:* Safe to drop rows or impute with mean/median.
- **MAR (Missing at Random):** The missingness is related to *other observed variables*. (e.g., Men are less likely to report their depression scores than women). *Solution:* Use predictive imputation (e.g., KNN Imputer) using the other variables.
- **MNAR (Missing Not at Random):** The missingness is related to the *unobserved data itself*. (e.g., People with very high incomes refuse to report their salary). *Solution:* Highly problematic. Requires domain expertise or specialized modeling.

5.4 Outlier Detection & Treatment

Outliers are extreme deviations from other observations. They can be genuine anomalies (a rare disease) or errors (a typo).

- **Z-Score Method:** Measures how many standard deviations a point is from the mean. $|Z| > 3$ is typically considered an outlier. Assumes normal distribution.
- **IQR (Interquartile Range) Method:** Robust to non-normal distributions.

$$IQR = Q3 - Q1 \tag{1}$$

Bounds are defined as $[Q1 - 1.5 \times IQR, Q3 + 1.5 \times IQR]$. Points outside these bounds are outliers.

Common Pitfall: Never Delete Outliers Blindly

In healthcare or finance, an outlier might be the most important data point! A sudden spike in server traffic could be a cyberattack; a rare symptom could be a novel disease. Always investigate outliers with domain experts before removing them.

6 The Python Data Science Toolstack

To execute the concepts learned in this course, we will use the industry-standard Python ecosystem.

6.1 The Core Libraries

- **Jupyter Notebooks:** Our interactive canvas. It allows us to combine executable code, rich text (Markdown), equations, and visualizations in a single document.
- **NumPy (Numerical Python):** The foundation of the ecosystem. It provides the `ndarray` object, enabling fast, vectorized mathematical operations on multi-dimensional arrays.
- **Pandas:** The Swiss Army Knife of data analysis. Built on top of NumPy, it introduces the `DataFrame` (a 2D tabular structure). It handles data I/O, cleaning, merging, and time-series manipulation.
- **Matplotlib & Seaborn:** The visualization engines. Matplotlib provides granular control over plots, while Seaborn offers beautiful, high-level statistical graphics.
- **Scikit-Learn:** The machine learning library. It provides unified APIs for preprocessing, model selection, supervised/unsupervised learning, and evaluation metrics.

Pro Tip: The "Tidy Data" Philosophy

As you transition into Pandas, adopt the "Tidy Data" framework:

1. Each variable forms a **column**.
2. Each observation forms a **row**.
3. Each type of observational unit forms a **table**.

Structuring your data this way from the beginning will save you hundreds of hours of frustration later.

7 Conclusion & Next Steps

You are embarking on a rigorous but highly rewarding journey. The demand for professionals who can translate raw data into strategic insights is growing exponentially across Africa and the globe.

7.1 Expectations for Success

- **Commitment:** This is an intensive 6-week program. Consistency is key. Code every day, even if it is just for 30 minutes.
- **Embrace the Struggle:** You will encounter errors, broken code, and confusing statistical concepts. This is normal. Debugging is where the actual learning happens.
- **Collaboration:** Data science is a team sport. Share your Jupyter notebooks, help your peers debug, and discuss your capstone ideas.

Welcome to the Data & AI ecosystem. Let's get to work!

Victor Wandera Lumumba, MSc.

Statistician — Data Analyst — ML Expert

lumumbavictor172@gmail.com • beyonddataanalytics.online